



Map organisation and temporal boundaries

V and AV-Speech share the same frozen visual representation. Speech changes the semantic map supplied to Direct C.

Human reference - schematic, not time-scaled



Model maps - a shared linear time axis (seconds)



V - 6 regions

- C1** 0-180 s
Car-door activity
- C2** 180-360 s
Walking / cars / phone
- C3** 360-540 s
Fence / residential approach
- C4** 540-720 s
Gun display / restraint
- C5** 720-1035 s
Ground control / monitoring
- C6** 1035-1235 s
Sustained limb control

AV-Speech: C1-C5

- C1** 0-180 s
Car-door + instruction
- C2** 180-360 s
Walk + armed context
- C3** 360-540 s
Fence + coordination
- C4** 540-585 s
Gun drawn
- C5** 585-675 s
Restraint + checking

AV-Speech: C6-C10

- C6** 675-765 s
Gloves + coordination
- C7** 765-900 s
Emergency response
- C8** 900-1035 s
Ground + commands
- C9** 1035-1215 s
Limb control + EMS
- C10** 1215-1235 s
Final hold / readiness

Region labels are condensed, not verbatim model outputs. Full Coarse summaries and uncertainty notes remain in the thesis appendix. The human reference is post-hoc and did not guide navigation or frame selection.